

Shared UI State and Host Enforcement: An Empirical Technical Note on Three Agent-UI Configurations

Eduardo Arana

Independent researcher

26 September 2026

Abstract

Under the tested configurations and one fixed write order, reusing an identifier overwrites in the shared json-render and A2UI mappings and is retained in separate MCP App instances. This empirical technical note audits an existing deterministic TypeScript harness across surface handoff, fixed-order identifier collision, and action round-trip, in three concrete agent-UI configurations. Node assertions establish the four-write outcome: replacement of the json-render summary and retention of both summaries in separate MCP instances. Every browser run instead exercises a separate two-write summary fixture; it corroborates the rendered outcome only for that fixture and does not validate the Node four-write scenario. Executed A2UI browser assertions show finance’s headline visible and risk’s absent, not an exact node count. The installed enforcing host handler rejects one model-only tool call; SDK-default `oncalltool` forwards it, a measured specification-versus-SDK gap. Interpretation must therefore name the adapter mapping, namespace/topology, and enforcement site together with the SDK or protocol. A claim-to-check mapping distinguishes supported, partial, and proposed claims; routing to every intended control owner remains partial, and a matched-topology intervention remains proposed, not performed. The evidence supports configuration-specific findings, not isolated protocol causality or rankings. No LLM experiment or performance measurement is reported. Code and evidence are archived with a persistent DOI.

1 Introduction

The existing harness exposes two configuration-dependent findings. Under the tested configurations and one fixed write order, reusing a summary identifier overwrites in the shared json-render and A2UI mappings and is retained in separate MCP App instances. Executed A2UI browser assertions confirm finance’s headline is visible and risk’s is absent in that surface. Separately, the tested MCP host rejects a model-only call with an enforcing handler but forwards it in the permissive variant. A comparison must identify the adapter mapping, namespace/topology, and enforcement site alongside the SDK or protocol: those choices are part of the observed system, not incidental implementation details. Figure 1 summarizes the adapter-recorded classifications for the three configurations.

This empirical technical note examines three concrete tested configurations, not three protocols in isolation. Protocol effects cannot be separated from adapter and topology here; a matched-topology comparison remains proposed, not performed. Its contribution is an auditable account of a small deterministic harness: research questions, a falsifiable hypothesis, bounded findings,

Capability (adapter classification)	json-render	A2UI	MCP Apps
Continue another agent's surface	caveat	supported	not expressible
Attribute a write to an agent	not expressible	not expressible	enforced
Isolate agents from each other	not expressible	caveat	enforced
Compose into one view	caveat	supported	not expressible
Retain both summaries after a fixed-order collision	write lost	write lost	separate instances
Route an action to its agent	out-of-band	out-of-band	supported
Gate a privileged action	caveat	not expressible	enforced

Figure 1: Adapter-recorded classifications per capability for the three tested configurations. Each cell is the label the adapter recorded while driving the installed SDKs through the scripted scenarios (one retained execution per scenario; generated from the recorded traces in `docs/COMPARISON.md`). Labels are classifications of these configurations, not independent measurements or protocol rankings (C8); the collision and host-enforcement rows are backed by separate assertions (C2, C5).

an evidence map, and a local reproduction procedure. UI descriptions, execution environments, and authorization decisions operate at different layers; collapsing those layers would misattribute namespace choices or host policies to a format. The retained comparison report is an artifact to audit, not independent scientific evidence. Scope and validity limits below delimit the findings; no exhaustive novelty claim is made.

2 Research questions

RQ1: Where do the existing adapters place the output of a scripted surface handoff?

RQ2: What survives a fixed-order identifier-collision probe in the configured namespaces? Under the tested configurations and one fixed write order, does reusing an identifier overwrite in the shared json-render and A2UI mappings while being retained in separate MCP App instances? This cannot generalize to alternative orders, scheduling or concurrency: only the one specified sequential write order is observed.

RQ3: In the tested paths, which local bookkeeping and host decisions determine action routing and one model-only tool rejection? In this MCP Apps topology, originating-agent identity comes from connection binding, not the payload.

3 Methodology

This retrospective case study examines deterministic scripted agents, with no LLM experiment. The unit of comparison is an SDK/adapter/host configuration. Even uniform scenario inputs could not isolate topology, adapter, or protocol causality. This is a bounded configuration study, not a protocol ranking or a test of security or composition tradeoffs. Browser MCP inputs are not identical to Node inputs. Three scenarios exercise three configurations, with an existing permissive-host variant for MCP Apps in S3.

3.1 A falsifiable local hypothesis

H1a (Node, four writes): json-render retains finance’s headline and MCP Apps retains both headlines in separate instances. H1b (browser, two writes): A2UI’s browser fixture displays finance’s headline and no risk headline within its surface; its Node check establishes only finance writer bookkeeping. These are independent fixture-specific predictions, not a common full-scenario browser validation. H1a and H1b observe only the one specified sequential write order and cannot generalize to alternative orders, scheduling or concurrency. A snapshot or DOM assertion showing the risk headline instead of finance in a shared summary, or loss of either isolated headline, would contradict this hypothesis for the tested configuration.

The S2 script awaits each emission: risk detail, finance detail, risk summary, then finance summary. This is a fixed-order identifier-collision probe, not actual concurrency or a scheduling benchmark. The focused check uses the existing snapshot assertions and browser assertions, rather than treating the adapter’s **LOST** or **ENFORCED** labels as measurements. Fresh execution status is recorded separately from this source-based hypothesis.

3.2 Configurations and versions

The manifest specifies ranges. The lockfile and locally installed package metadata agree on the exact versions below. Package versions are not protocol schema versions. These are observations of this checkout, not recommendations or statements about the newest available releases.

Configuration	Manifest range	Lockfile and installed version
json-render core and React renderer	~0.21.0 each	0.21.0 each
A2UI web core and Lit renderer	~0.11.0 each	0.11.0 each
MCP Apps extension SDK	~2.0.0	2.0.0
MCP client, core, and server	~2.0.0 each	2.0.0 each
React and React DOM	~19.2.3 each	19.3.0 each
Lit	~3.3.3	3.3.3
Zod	~4.2.0	4.6.5

The A2UI adapter imports `@a2ui/web_core/v0_9` and emits `version: "v0.9"`. Its package version 0.11.0 does not mean that the experiment uses a v0.11 protocol. MCP Apps SDK 2.0.0 likewise does not establish a wire protocol version. The extension specification and SDK are separate references; no exact upstream specification commit or negotiated wire-version value was retained for this paper. Live repository documentation was inspected on 21 September 2026 and may differ from release-time documentation. The local lockfile fixes the tested package resolution, not the mutable web references.

The json-render adapter maintains one **Spec** per conceptual surface and maps logical block, row, and control IDs to element keys. It calls **applySpecPatch** for element additions/replacements and also directly updates root children. Therefore, its patch log does not capture every mutation. Its catalog validates emitted specs; **writtenBy** is an ordinary catalog prop and **lastWriter** is adapter bookkeeping. Neither is authenticated provenance.

The A2UI adapter validates messages using **A2uiMessageSchema** and drives **MessageProcessor**. Its component IDs and data paths derive from block IDs. A shared **rootChildren** collection coordinates the surface root. Node uses a schema-only harness catalog; the browser uses the Lit basic catalog. Snapshot titles and writer names partly come from adapter maps, while row values come from the data model. Surface addressing is not authorization: an address identifies a target but

does not establish a caller’s permission.

The MCP Apps adapter creates one server per agent and one app instance per surface/agent pair. Its Node path uses real SDK clients, servers, apps, and bridges connected by in-memory transports. Its snapshot is reconstructed from adapter-held blocks, not from a browser DOM. The browser host mounts an iframe with `sandbox="allow-scripts"` for each agent and uses `PostMessageTransport`. This topology is an implementation choice. Connection identity is not universally agent identity: the mapping works here because the host establishes one server per named agent. A server could aggregate multiple agents in a different application.

Figure 2 locates the shared state, local ownership mapping, and enforcing handler in these configurations. The dashed path denotes the existing permissive variant, not a second enforcement mechanism.

3.3 Scenario design and evidence collection

S1 opens a trip surface, emits a planner itinerary, hands off to booking, and emits a reservation. S2 performs the fixed-order writes specified in H1. S3 emits basket and payment controls and activates ordinary and privileged actions. The amounts and booking operations are fixtures: the MCP server records calls and returns text; it does not charge money. The three scenarios are fixed examples, not a sample drawn from a population of agent workflows.

The runner constructs fresh adapters per scenario and configuration. The orchestrator is a scripted stand-in with named agents and a control-to-agent routing table, not an autonomous model-driven system. We inspect the existing tests and rerun them where local prerequisites permit; no scenarios are added or their operations or write order are modified. MCP’s collision classification records structural isolation from this adapter/topology separately from the tested host tool-visibility handler; it does not establish a host collision-enforcement policy.

The report generator takes adapter-authored classifications and uses first-recording-wins selection for each capability and configuration. It does not independently measure the meaning of `SUPPORTED`, `SUPPORTED_WITH_CAVEAT`, `LOST`, `ENFORCED`, or `NOT_EXPRESSIBLE`. Repeated labels do not constitute independent observations. Timestamp normalization makes its Markdown stable but does not validate its interpretation. A freshness check establishes correspondence between report and generator, not the truth of every sentence in the report.

We therefore distinguish classifications from assertions on snapshots, messages, tool responses/refusals, and rendered elements. The conformance tests validate the S1 A2UI messages and json-render catalog output, examine recorded patches, and check report and image-link freshness. They do not establish full specification conformance. The A2UI deletion message is not exercised in that S1 check. The json-render patch assertions do not cover its direct root mutation.

Specifically, `Tracer.record` stores the outcome supplied by the adapter; `outcomeOf` returns the first matching entry. In scenario tests, json-render’s S3 confirmation check and MCP’s permissive S3 check assert these labels only. MCP’s S2 `SUPPORTED_WITH_CAVEAT`/`NOT_EXPRESSIBLE` labels do not measure collision prevention or impossibility of composition; the former records structural isolation in the adapter/topology, not a tested host collision policy. Separate assertions check the two retained headline values. A2UI S2 checks labels and the writer map, not Node-rendered text. The bridge tests instead exercise calls and assert actual returned content or rejection plus refusal records. None of these labels is an empirical score.

Browser evidence has an additional qualification: the MCP browser host contains its own fixture script. In S2 it renders the two summaries without the two uncontested detail blocks used by the Node scenario: Node awaits risk-detail, finance-detail, risk/summary, finance/summary; browser




Figure 2: Configured state and enforcement boundaries. Panel A summarizes two separate declarative-renderer runs; panel B shows the app-to-server call path for each MCP instance. Writer maps are bookkeeping. In the tested enforcing variant, the installed host handler rejects the model-only call; the permissive variant forwards it. Original source-derived illustration, not a fresh browser observation.

MCP awaits risk/summary, finance/summary only. Its DOM assertions establish those browser-specific summary values, not validation of the same full Node S2 scenario. Retained screenshots alone cannot establish that a fresh browser run passed.

4 Bounded observed results

4.1 RQ1: Handoff and placement

Existing S1 assertions check a single `trip` snapshot region containing both `itinerary` and `reservation` in `json-render` and A2UI. MCP Apps checks separate `trip::planner` and `trip::booking` regions. The browser assertions check both blocks within one React tree or A2UI surface, and separate MCP iframes. These checks support a placement observation under the chosen mappings. They do not prove that MCP Apps cannot represent an aggregated UI or that either declarative format provides authenticated cross-agent ownership.

4.2 RQ2: Fixed-order identifier collision

These results observe only the one specified sequential write order and cannot generalize to alternative orders, scheduling or concurrency. The current-source browser run verifies the fixture; earlier screenshots are retained as historical artifacts.

The focused local S2 rerun passed its three selected tests; ten non-S2 tests in that file were skipped by the filter. The `json-render` assertion checks the finance writer and headline, plus survival of the two uncontested blocks. The A2UI Node assertion checks the finance writer from adapter bookkeeping; its browser assertion checks headline text. MCP snapshot assertions check both headlines in different instances. H1 has fixture-specific support from these Node checks and the canonical current-source browser run `browser-s2-fixed-order-cFIGKT5C/` on 25 September 2026 UTC. A2UI's browser assertion checks finance's headline visible and risk's headline absent within the surface. It does not assert an exact summary node count, despite its test title. The execution ledger identifies a raw browser report for the current-source run as recording 18 passed tests: nine render assertions and nine screenshot tests. The complete current browser run resolves to immutable `browser-s2-fixed-order-cFIGKT5C/` (25 September 2026 UTC; commit 378aabee; Node v23.5.0) and includes that report, source snapshots, build hashes, and capture hashes. The MCP browser fixture omits Node's risk-detail and finance-detail writes. Its two summary assertions are a separate browser result, not corroboration of the full four-write Node S2 or support across a common shared fixture.

Under the tested configurations and one fixed write order, reusing an identifier overwrites in the shared `json-render` and A2UI mappings and is retained in separate MCP App instances. The replacement is consistent with reusing a slot in these shared namespaces. It is not an observed scheduling race. The inspected adapter records a changed writer; this trace is not an audit of all conflict signals available in a protocol or SDK. Namespacing, ownership checks, arbitration, or shared aggregation could change the outcome; none is evaluated as an intervention here.

Figure 3 separates the four awaited writes from the retained summary states. The arrows express script order, not elapsed time or competing schedules.

4.3 RQ3: Routing and host decisions

In the shared-state configurations, ordinary action routing uses the orchestrator's control table. Tests inspect routing entries and deliveries, and the A2UI action envelope carries surface/component fields and a timestamp without an agent field. These payload observations do not rule out an application-supplied identity envelope. Encoding a name in context or props would still require trusted binding before it could authorize anything.

`json-render`'s privileged control includes `confirm` metadata. The Node adapter constructs an




Figure 3: S2 fixed-order collision: the four-write sequence is the Node fixture. Node assertions check the json-render headline and both MCP headlines. Executed A2UI DOM assertions check finance’s headline visible and risk’s absent (C3, supported), not an exact node count. The separate MCP browser fixture executes only risk/summary then finance/summary, omitting both detail writes; it does not validate full Node S2. The two MCP summaries belong to separate instances, not a merged surface. Arrows indicate fixed script order, not time or scheduling.

event and classifies that metadata; it does not run an independent authorization service. The browser suite clicks the ordinary json-render action, not the privileged confirmation flow. Confirmation UI is not independent authorization, and the existing assertions do not establish complete consent behavior or protection against an omitted confirmation binding. A2UI has no privileged gate implemented in this adapter; that is not a proof that its hosts or trusted catalog components cannot impose one.

For MCP Apps, the host installs one enforcing `oncalltool` handler after `bridge.connect`. The handler looks up the tool and rejects one model-only call. The bridge tests check successful app-visible calls, rejection messages and refusal records, and successful model-only calls when enforcement is disabled. The permissive host variant forwards that call. This is a measured specification-versus-SDK gap for the installed enforcing host handler and SDK-default `oncalltool`, not a security guarantee or evidence about every MCP host or SDK release. R4 reports the external premise that the inspected extension specification assigns visibility enforcement to the host, but that premise is unverified within this artifact.

Routing evidence is narrower than some trace labels suggest. For actions, the Node MCP adapter selects the first instance for a conceptual surface. Its helper is `findInstanceForSurface`. It does not find the agent owning the clicked control. The routing test only requires some delivery without the fallback table. Thus it cannot prove correct owner routing for every S3 control. Browser S3 instead clicks inside the cart and payments frames separately. We retain this distinction and leave

the existing behavior unchanged.

5 Claim-evidence mapping

Here **supported** means a bounded statement has a matching inspected assertion or direct source observation, with run status separately recorded. **Partial** means a relevant assertion exists but leaves a material inference untested. **Proposed** means a future check has not been performed. These terms are this paper’s evidence statuses, distinct from the adapter outcome vocabulary.

Claim	Status	Evidence and falsifier
C1: S1 yields one shared region or two MCP instances in this topology.	Supported	S1 snapshot region and block assertions; contradicted by missing blocks or different region counts.
C2 / H1a (Node, four writes): Under the tested configurations and one fixed write order, reusing an identifier overwrites in the shared json-render and A2UI mappings and is retained in separate MCP App instances.	Supported	Node headline assertions; contradicted by different retained values. The browser fixture is not evidence for H1a.
C3 / H1b (browser, two writes): Under the tested A2UI shared mapping and fixed write order, the fixture displays finance’s headline after the later write.	Supported	The cited current browser run resolves to immutable <code>browser-s2-fixed-order-cFIGKT5C/</code> (25 September 2026 UTC; commit <code>378aabee</code> ; Node <code>v23.5.0</code>) and records assertions checking finance visible and risk absent within the surface. It is distinct from H1a and does not assert an exact node count.
C4: S2 is sequential, with no scheduling test.	Supported	Four awaited emissions in the S2 source; contradicted by an overlapping scheduler in the executed path.
C5: The enforcing MCP handler rejects the model-only call, while the permissive variant forwards it.	Supported	Bridge rejection/refusal and successful-result assertions; contradicted by reversed or identical behavior.
C6: Every MCP S3 control reaches its intended agent.	Partial	Existing test checks some table-free delivery; first-instance selection prevents this stronger inference.
C7: json-render includes confirmation metadata, with consent enforcement incompletely checked.	Partial	Adapter source includes metadata; the scenario test asserts its classification only, not independent policy or privileged browser consent.
C8: Matrix labels are adapter classifications, not independent measurements.	Supported	Source shows adapter-authored classifications; freshness checks correspondence only.

Claim	Status	Evidence and falsifier
C9: A matched-topology intervention isolates protocol causality.	Proposed	Pending controlled comparison; no new experiment is reported. The one-server-per-agent topology has not been ablated, so no collision, isolation, or host-policy result is protocol-level. All collision findings are conditioned on the one fully deterministic write order and cannot be extrapolated to concurrent or out-of-order writes.

Each **Supported** row reflects exactly one retained execution (n=1) identified for its check; the harness uses deterministic scripted agents rather than an LLM experiment. No mean, variance, pass rate, or run-to-run browser stability is claimed.

The linked scenario, bridge, conformance, and browser test files in Methodology and Results are the check locations for C1–C8. This mapping exposes uneven coverage rather than assigning a numeric score across unlike checks.

For C1–C3 and the tested browser actions, the renamed render spec supplies the assertions and the screenshot spec supplies nine captures. By contrast, the earlier pre-`browser-terminal-20260922T013449Z-zH4HTK/` records are historical only and are not claim-support evidence for this fixture. The execution ledger records the single cited current browser run, which resolves to immutable `browser-s2-fixed-order-cFIGKT5C/` (25 September 2026 UTC; commit 378aabee; Node v23.5.0). Its report, source snapshots, build hashes, and capture hashes bind the current browser claim for this fixture. Other retained browser directories have different source snapshots and are execution records, not redundant confirmation of this claim. This does not make the distinct MCP browser fixture corroborate the four-write Node fixture. C6 remains partial: successful browser actions in its distinct fixture do not repair the Node first-instance routing gap.

6 Related work

The json-render repository describes catalog-defined components and actions rendered from structured specifications [R1]. A2UI similarly separates declarative UI data from client implementation and describes a trusted component catalog [R2]. These are meaningful security properties: constraining requested components and handlers can limit what untrusted UI data asks the client to run. They do not by themselves establish per-writer authorization in our shared namespace. Conversely, our collision example is not evidence of blanket insecurity in declarative rendering.

The MCP Apps repository describes tool-associated UI resources and a host bridge [R3]. Its specification describes visibility checks and sandbox/CSP obligations [R4]. Those obligations are not proof that this harness checks all of them. In particular, iframe sandbox attributes and successful bridge tests do not establish exhaustive CSP enforcement, transport security, or resistance to malicious content. Declarative catalogs and sandboxed app delivery are compatible layers: a hosted application could use a declarative renderer internally. No integration experiment is reported here.

The claim-status vocabulary is defined in this paper’s Claim-evidence mapping: it separates an inspected assertion or direct source observation from a material untested inference and from a proposed future check. It does not depend on an unavailable external methodological source. This

discussion is selective, not an exhaustive literature or novelty review.

7 Threats to validity

Internal validity. Identical logical inputs leave adapter bookkeeping, shared-root coordination, naming conventions, host enforcement, catalogs, rendering frameworks, transports, and one-server-per-agent MCP topology as confounders. The default/permissive MCP comparison narrows one host-policy question, but it is not a factorial comparison across all configurations. Accordingly, the isolation and collision findings describe the tested adapter and topology configurations, not protocol-level enforcement results; a matched-topology intervention remains unperformed.

Construct validity. A region count is a property of the adapter snapshot; it is not a usability measure. A writer prop is not authenticated provenance. Addressability is not authority. A confirmation binding is not an authorization decision. Adapter labels and some scenario assertions share their interpretation, creating a risk of circular validation. The editorial and manuscript regression tests are likewise authored and iteratively adjusted by the same AI-assisted process that produced the prose they check. The independent aspects of the checks are the inspected contents, response/refusal assertions, and DOM expectations; they are still authored within this repository, not by independent evaluators.

Execution and coverage. Node transports do not test browser isolation. The MCP browser fixtures differ from the Node scenario. Historical screenshots remain retained artifacts; the nine new captures are separately archived and do not establish historical reproduction. Four post-rename attempts are retained: two completed with 18 passing tests, and two aborted before any test body ran. The two completed executions used different source revisions: their snapshots differ for `src/adapters/mcp-apps/adapter.ts`, `src/orchestrator/types.ts`, and `src/report/generate.ts`. Each reported pass count is from one execution ($n=1$), not repeated trials; no mean, variance, or stability statistic is claimed. No per-test timing or other run-level equivalence is claimed. No completed execution was discarded, and the immutable current-source run is `paper/evidence/browser-s2-fixed-order-cFIGKT5C/` (25 September 2026 UTC; commit `378aabee`; Node `v23.5.0`). Neither completed run is an independent reproduction: both ran on the same machine, from the same checkout, in the same session. No adversarial, forged-identity, CSP, or cross-server test was run. The suite also does not comprehensively test transport authentication, authorization revocation, or arbitrary concurrent schedules. No completed security review is claimed.

Classification scope. The matrix contains many individually classified configuration/scenario cells. Because these are deterministic code-behavior observations rather than independent statistical tests, no multiple-comparison estimate is reported. The number of cells nevertheless increases the surface for a stale or incorrect classification; the freshness check tests report/code correspondence, not the truth of every classification.

External and temporal validity. Three designed scenarios, fixed fixture data, and the pinned installed packages cannot establish population effects or general protocol superiority. Later SDKs, other renderers, different server topologies, and stricter application policies can behave differently. Mutable upstream documentation is contextual evidence, not a version-pinned release archive. No latency, throughput, cost, user study, or model capability was measured.

8 Data availability

The exact evidence reported here is archived as v0.2.0 at <https://doi.org/10.5281/zenodo.22965404>. The concept DOI, <https://doi.org/10.5281/zenodo.22896881>, resolves to the latest version. The archive contains the harness, the editable manuscript, a frozen v0.1 Markdown snapshot, and the versioned evidence records cited in this paper. The harness code is released under the Apache License 2.0; the manuscript, figures and evidence records under CC BY 4.0, as identified by the root `LICENSE` and package metadata.

9 Reproducibility

The reproduction guide gives commands and expected semantic outcomes, and the execution record lists every command run, including failed attempts and unavailable prerequisites, with retained artifacts. Inputs, tooling, and evaluation configuration are versioned now; commit-level provenance before 22 September 2026 does not identify the complete paper workspace. Separate paper checks and SHA-256 manifests record the manuscript, tooling, figures, and generated outputs.

Semantic reproducibility means obtaining the specified regions, values, events, and refusal behavior under the same configuration. Byte identity is a separate question. A2UI action timestamps vary; report generation normalizes those timestamps. We check generated LaTeX against fresh Pandoc output under the recorded toolchain, not PDF byte identity across machines. The Markdown manuscript is the only editable paper source; LaTeX and PDF are generated derivatives.

The Node test suite (24 tests across three files) and TypeScript checking pass. The browser claims rest on the current browser run, which resolves to immutable `browser-s2-fixed-order-cFIGKT5C/` (25 September 2026 UTC; commit `378aabee42a69c61edc7d7a37c934465b4a66e30`; Node `v23.5.0`) and records current-source and build hashes, 18 passing tests, and nine capture hashes. All runs were performed on a single machine by the author; independent reproduction has not been performed.

10 Conclusion

S2 observes only the one specified sequential write order and cannot generalize to alternative orders, scheduling or concurrency. The retained pre-rename browser findings are historical; the complete current browser run, resolving to immutable `browser-s2-fixed-order-cFIGKT5C/` (25 September 2026 UTC; commit `378aabee`; Node `v23.5.0`), supplies current browser-source provenance for the renamed fixture.

Across three concrete tested configurations, handoff placement and retained values reflect the chosen shared or separated mappings. Under the tested configurations and one fixed write order, reusing an identifier overwrites in the shared json-render and A2UI mappings and is retained in separate MCP App instances. Node assertions establish the four-write S2 outcome for those configurations. Browser MCP executes only risk/summary then finance/summary. Those browser observations corroborate rendered outcomes only for their own fixtures and cannot validate the Node four-write fixture. Neither fixture tests concurrency, scheduling, interleaving, or randomization. MCP model-only tool rejection is measured behavior of one installed enforcing host handler, whereas SDK-default `oncalltool` forwards it; this is a specification-versus-SDK gap, not a security guarantee. In the tested topology, originating-agent identity comes from connection binding rather than the payload, while correct routing for every control remains only partially checked. Protocol effects cannot be separated from adapter and topology here; a matched-topology comparison remains

proposed, not performed. These results do not isolate protocol causality, establish protocol-level collision enforcement, or establish unavoidable incompatibility between composition and security. No adversarial, forged-identity, CSP, or cross-server test was run. Integrations remain possible, and trusted catalogs retain meaningful security properties.

Future experiments are explicitly pending: matched namespaces and topology, controlled scheduling with overlapping writes, trusted identity envelopes, privileged confirmation interaction tests, and complete host-policy tests. They are suggestions for separate approved work, not results of this paper.

11 AI-assistance disclosure

An OpenAI Codex coding assistant inspected local source and accessible references, assisted with drafting and editorial revision, prepared source-derived figures and supporting documentation, added focused manuscript checks, and executed the recorded local validation commands. Anthropic Claude assisted with editorial revision. The author reviewed all claims, citations, limitations, and the final presentation, and takes full responsibility for the content. AI-assisted preparation is separate from the deterministic harness: no LLM participates in its scenario runs.

12 References

R1. Vercel Labs. `json-render` repository and README. Inspected 21 September 2026 at commit `3ad381881194e7011ad3ccd6d668033495a06c29`. Local SDK context: `@json-render/core` and `@json-render/react` 0.21.0. Used for catalog/renderer design, not a security proof.

R2. A2UI project. A2UI repository and README (the inspected `google/A2UI` URL redirects here). Inspected 21 September 2026 at commit `c08702a4bf8ab22a862cf5de8b5b1542be6081f6`. Local context: web core and Lit 0.11.0, and the adapter-selected v0.9 schema.

R3. Model Context Protocol contributors. MCP Apps SDK repository and README. Inspected 21 September 2026 at commit `6d9bdc7babf275b759225aa722cbf5510c4c6021`, the same revision R4 pins. Local context: extension, client, core and server packages 2.0.0. SDK package version is distinct from wire version.

R4. Model Context Protocol contributors. MCP Apps extension specification, `draft/apps.mdx`. Visibility and sandbox/CSP sections inspected 21 September 2026; source path pinned to commit `6d9bdc7babf275b759225aa722cbf5510c4c6021`. The pin was added after earlier retained evidence used a mutable main path; this document records the pin but does not include an archived copy or independent verification of it. Its reported host-requirement premise is therefore unverified within this artifact and is not evidence for the host-enforcement conclusion. The filename’s “draft” is retained as the source location, without claiming a release status from the SDK version.